

Energy Scope report

Date of the report: 2022/04/19 18:36:33

GENERAL INFORMATION

- Jobid: 20220419193252
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_19_04/159_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/19 19:32:52.656749
- Duration (including ES prologue and epilogue): 3803 (sec)


ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 1343.5 (J/sec)
- Application energy consumption measurement: 5109328 (J) 1.4193 (kWh)
- Global application energy consumption estimation: 6858708 (J) 1.9052 (kWh)
- Global application carbon production estimation (FR): 97.346 (gCO2)
- Energy efficiency (ref TDP): 59.18 (%)

Eprofile:



ENERGY ACQUISITION INFORMATION

- Period(ms): 1590.549
- Acquisition quality (low, medium, high): low
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	377003	63.2	59.1
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	364403	60.3	55.5
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	513057	54.0	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	531378	55.9	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	475803	50.0	20.0

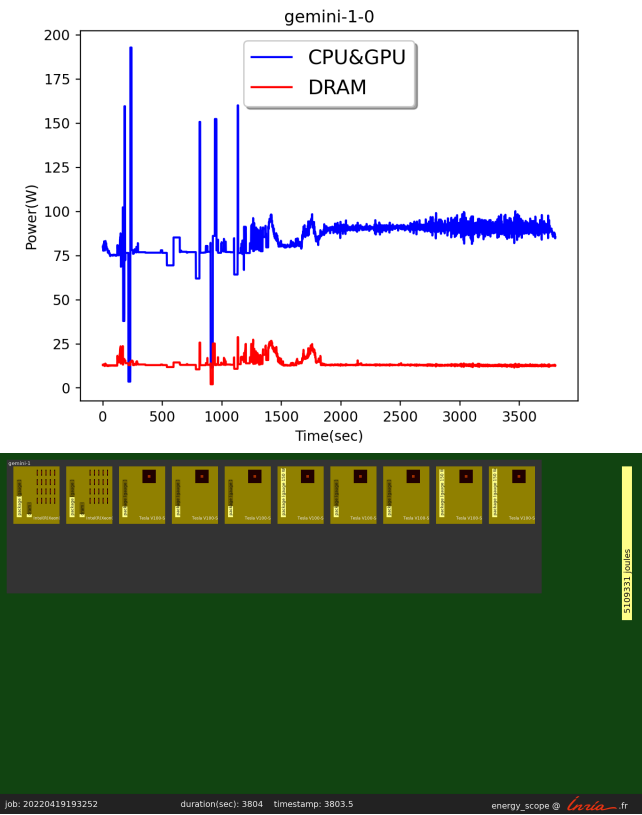
node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	605214	63.7	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	511142	53.8	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	524399	55.2	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	602326	63.4	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	604603	63.6	20.0

PROFILES and CORE TEMPERATURE


Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.


*node gemini-1/0




*node gemini-1/1




*node gemini-1/gpu-nvidia-0

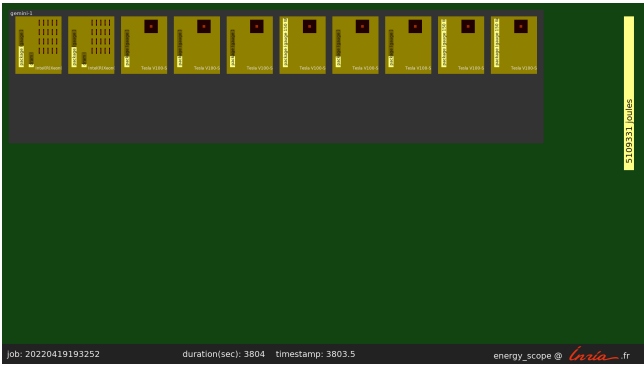


*node gemini-1/gpu-nvidia-1

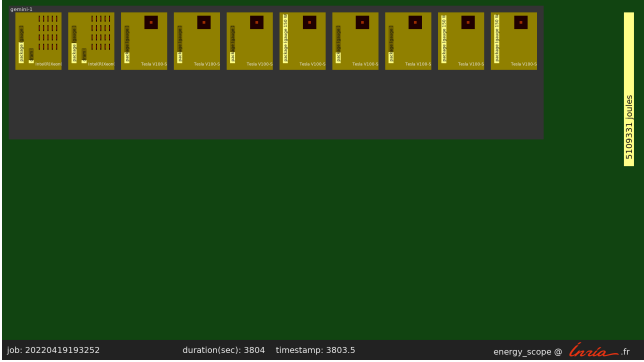



*node gemini-1/gpu-nvidia-2








*node gemini-1/gpu-nvidia-3




*node gemini-1/gpu-nvidia-4



*node gemini-1/gpu-nvidia-5



*node gemini-1/gpu-nvidia-6



*node gemini-1/gpu-nvidia-7

