

Energy Scope report

Date of the report: 2022/04/19 18:15:54

GENERAL INFORMATION

- Jobid: 20220419192634
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_19_04/543_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/19 19:26:46.331385
- Duration (including ES prologue and epilogue): 2931 (sec)

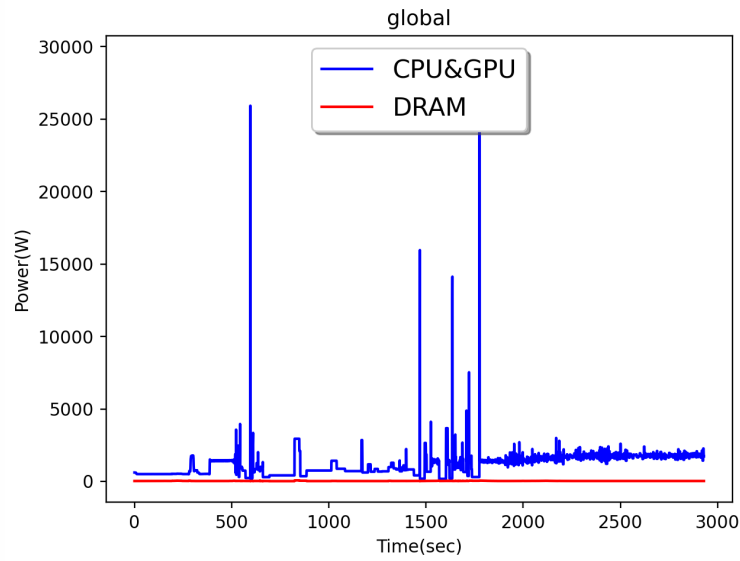
ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 1199.3 (J/sec)
- Application energy consumption measurement: 3515013 (J) 0.9764 (kWh)
- Global application energy consumption estimation: 4863273 (J) 1.3509 (kWh)
- Global application carbon production estimation (FR): 69.036 (gCO2)
- Energy efficiency (ref TDP): 52.83 (%)

Eprofile:



ENERGY ACQUISITION INFORMATION

- Period(ms): 1565.798
- Acquisition quality (low, medium, high): low
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	281797	60.7	58.9
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	272906	58.0	54.9
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	362893	49.5	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	370940	50.6	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	362301	49.4	20.0


node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	375642	51.3	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	375757	51.3	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	364988	49.8	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	372384	50.8	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	375405	51.2	20.0

PROFILES and CORE TEMPERATURE


Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.


*node gemini-1/0




*node gemini-1/1




*node gemini-1/gpu-nvidia-0



*node gemini-1/gpu-nvidia-1




*node gemini-1/gpu-nvidia-2

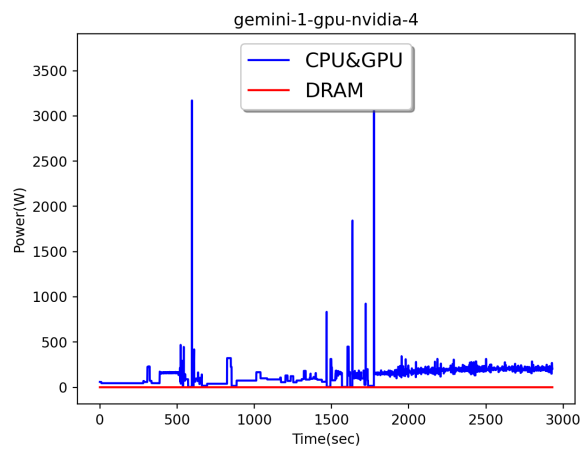





*node gemini-1/gpu-nvidia-3




*node gemini-1/gpu-nvidia-4



*node gemini-1/gpu-nvidia-5



*node gemini-1/gpu-nvidia-6



*node gemini-1/gpu-nvidia-7

