Energy Scope report

Date of the report: 2022/04/19 18:36:47

GENERAL INFORMATION

- Jobid: 20220419193246
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_19_04/688_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/19 19:32:46.583541
- Duration (including ES prologue and epilogue): 3815 (sec)

ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 1341.8 (J/sec)
- Application energy consumption measurement: 5118996 (J) 1.4219 (kWh)
- Global application energy consumption estimation: 6873896 (J) 1.9094 (kWh)
- Global application carbon production estimation (FR): 97.55 (gCO2)
- Energy efficiency (ref TDP): 59.11 (%)

Eprofile:

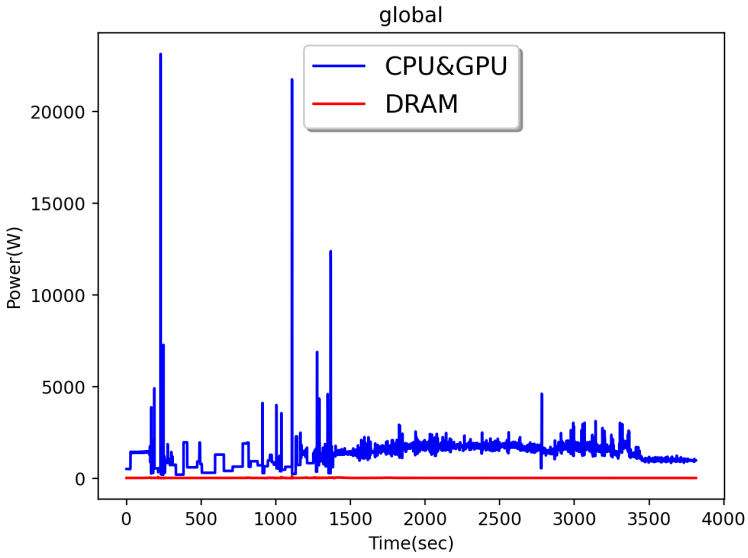

ENERGY ACQUISITION INFORMATION

- Period(ms): 1570.813
- Acquisition quality (low, medium, high): low
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	378189	63.2	59.1
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	365538	60.3	55.5
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	513606	53.9	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	531945	55.8	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	476360	49.9	20.0

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	606749	63.6	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	511698	53.7	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	524962	55.0	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	603838	63.3	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	606111	63.6	20.0

PROFILES and CORE TEMPERATURE

Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.

*node gemini-1/0

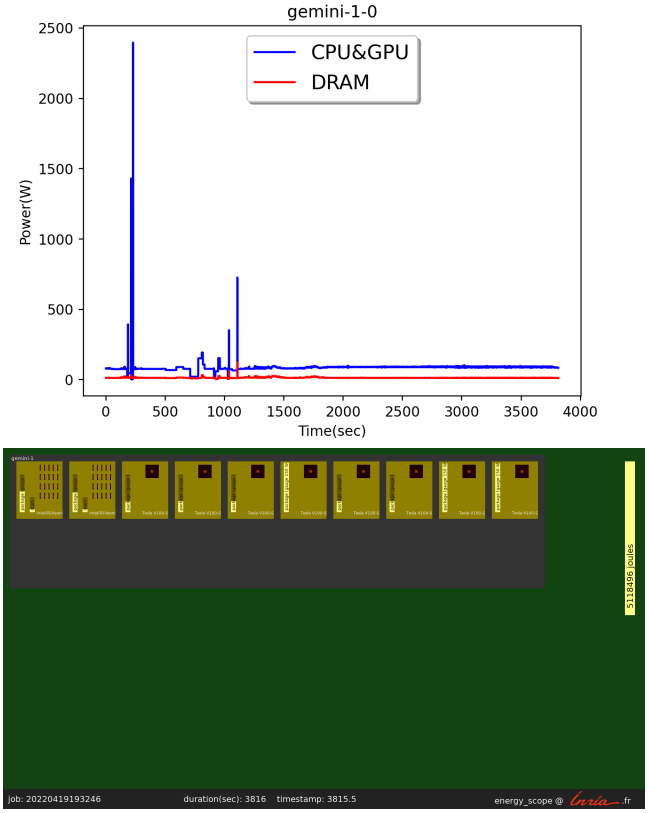

*node gemini-1/1

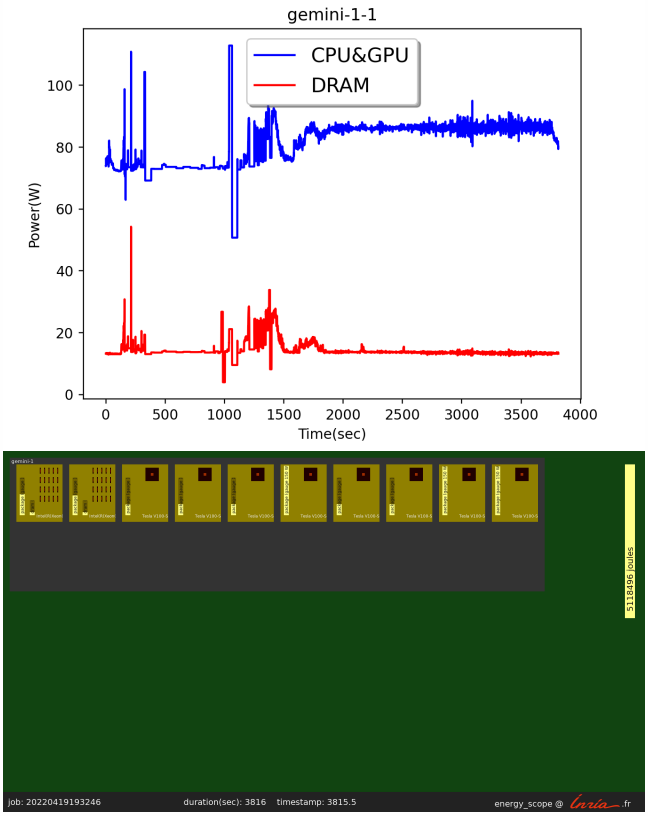

*node gemini-1/gpu-nvidia-0

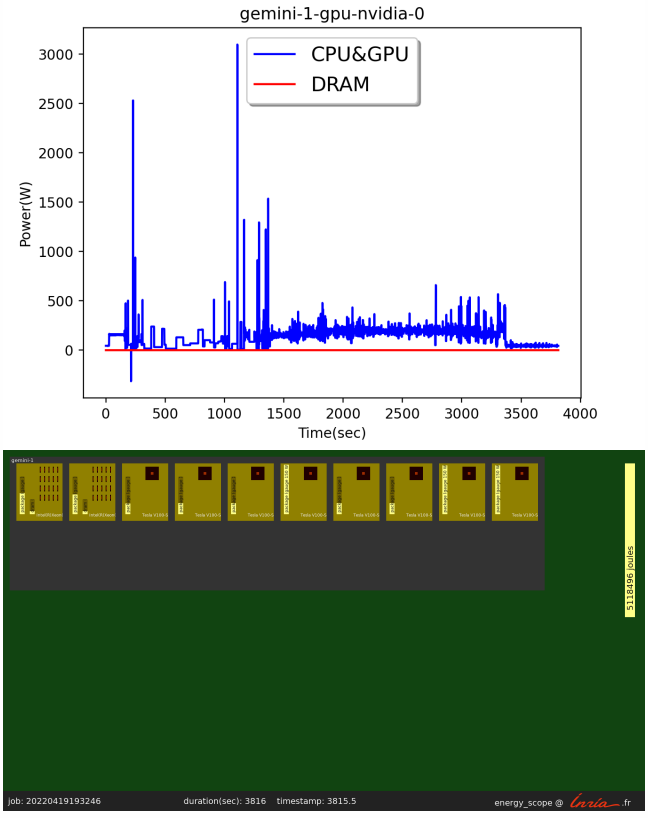

*node gemini-1/gpu-nvidia-1

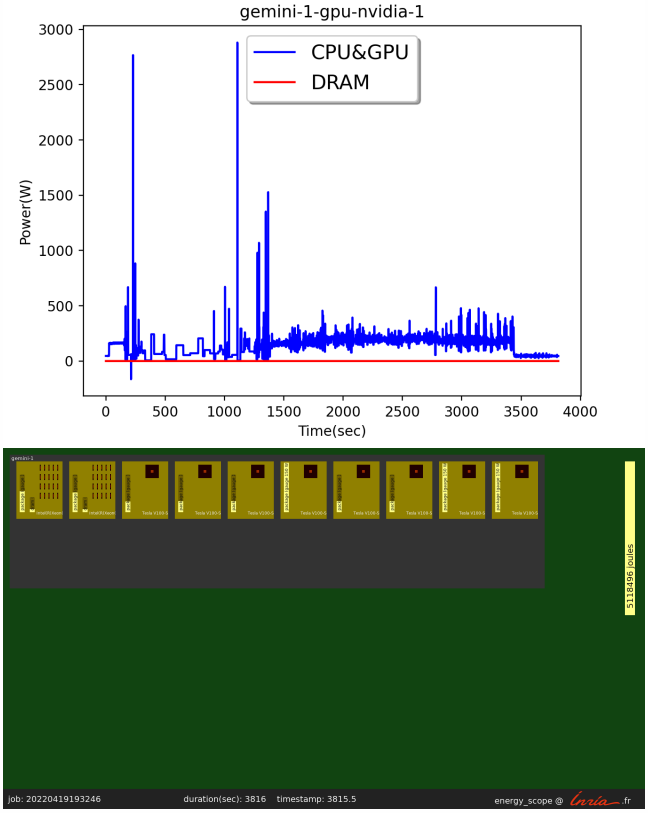

*node gemini-1/gpu-nvidia-2

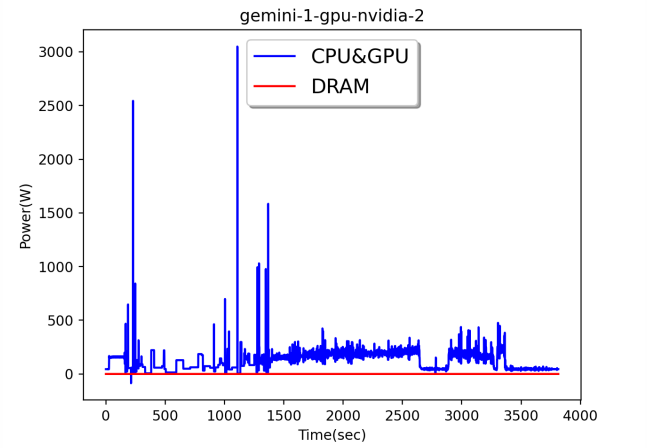

*node gemini-1/gpu-nvidia-3

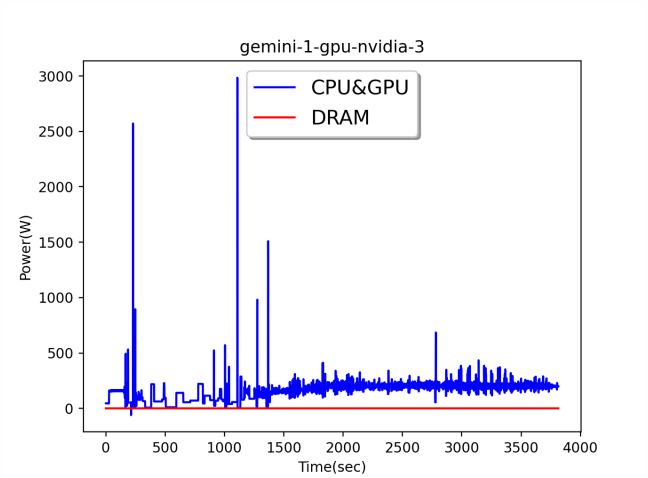

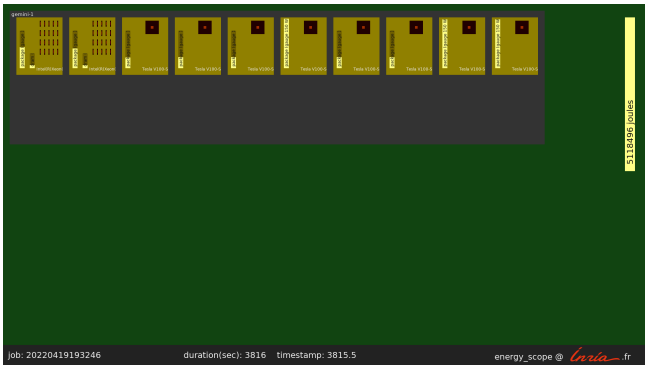

*node gemini-1/gpu-nvidia-4

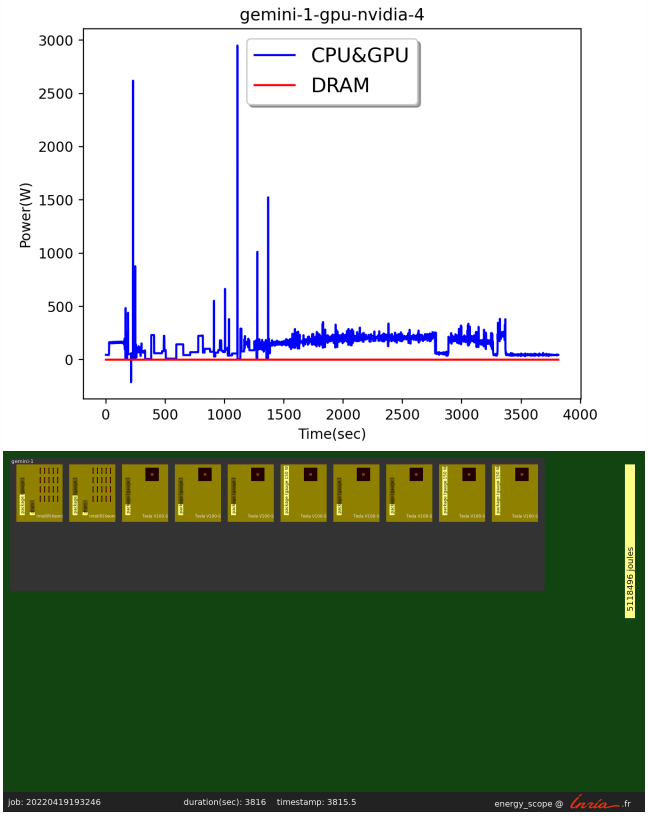

*node gemini-1/gpu-nvidia-5

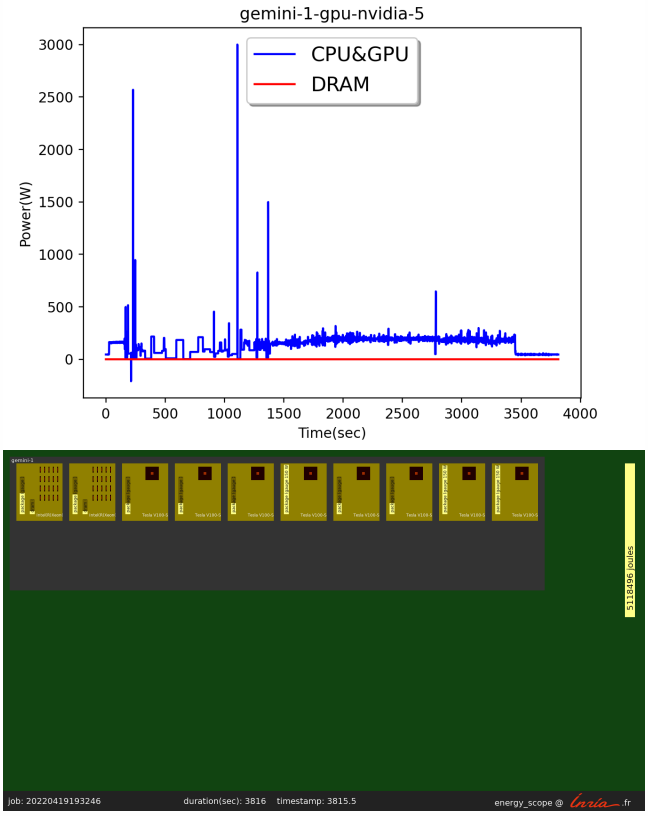

*node gemini-1/gpu-nvidia-6

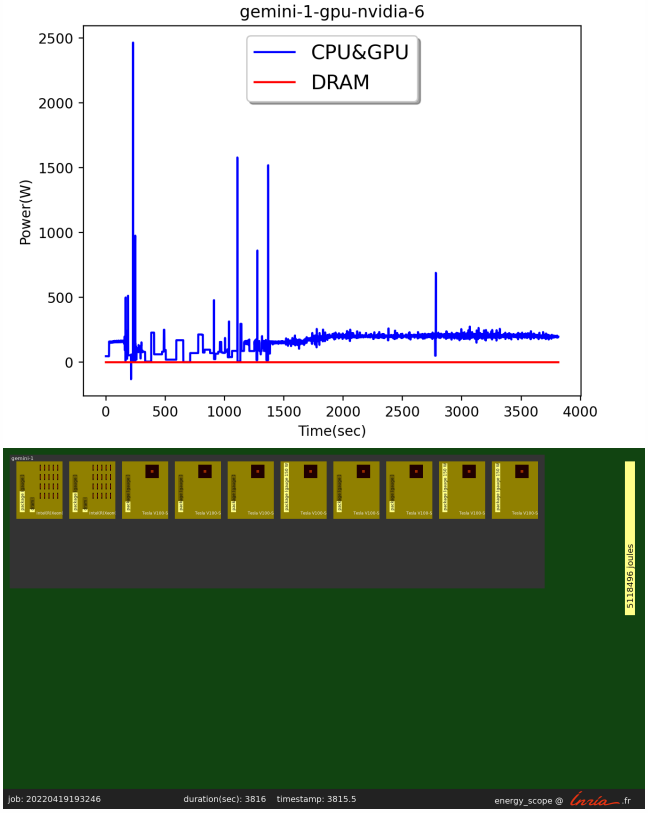

*node gemini-1/gpu-nvidia-7

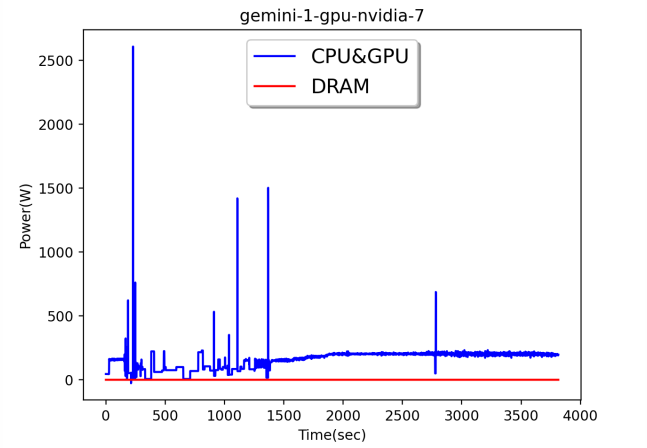


gemini-1
5118496 joules
job: 20220419193246
duration(sec): 3816
timestamp: 3815.5
energy_scope @ inria.fr