Energy Scope report

Date of the report: 2022/04/19 18:41:46

GENERAL INFORMATION

- Jobid: 20220419193321
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_19_04/760_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/19 19:33:21.941010
- Duration (including ES prologue and epilogue): 4086 (sec)

ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 1335.2 (J/sec)
- Application energy consumption measurement: 5455607 (J) 1.5154 (kWh)
- Global application energy consumption estimation: 7335167 (J) 2.0375 (kWh)
- Global application carbon production estimation (FR): 104.142 (gCO2)
- Energy efficiency (ref TDP): 58.82 (%)

Eprofile:

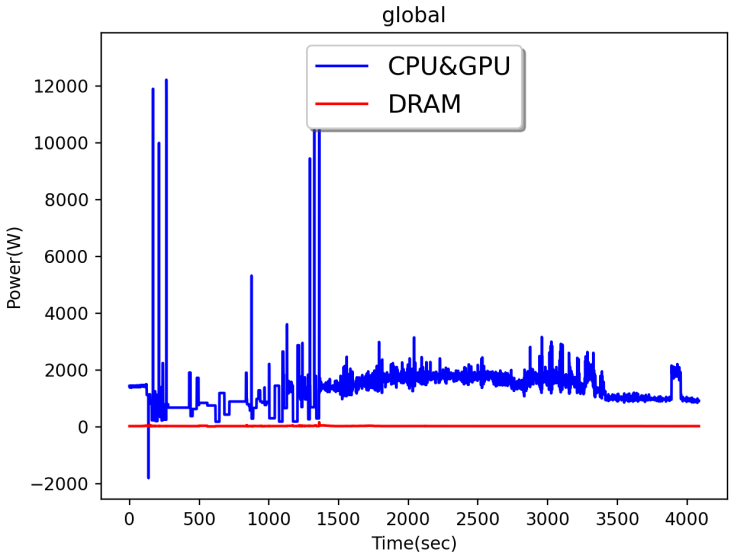

ENERGY ACQUISITION INFORMATION

- Period(ms): 1390.323
- Acquisition quality (low, medium, high): low
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	402869	62.9	58.8
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	389729	60.1	55.3
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	536822	52.6	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	555714	54.4	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	499972	48.9	20.0

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	663713	65.0	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	535366	52.4	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	548707	53.7	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	659246	64.5	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	663469	65.0	20.0

PROFILES and CORE TEMPERATURE

Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.

*node gemini-1/0

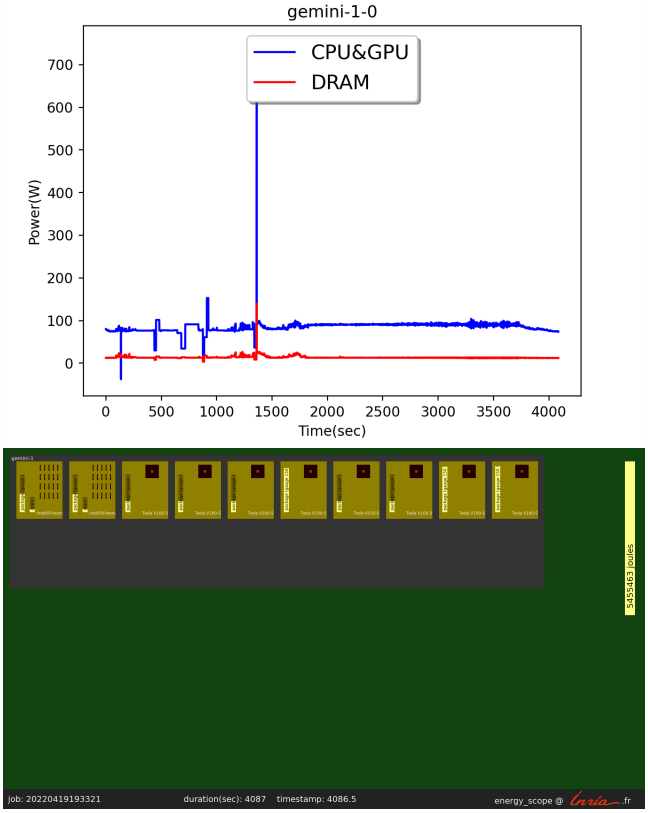

*node gemini-1/1

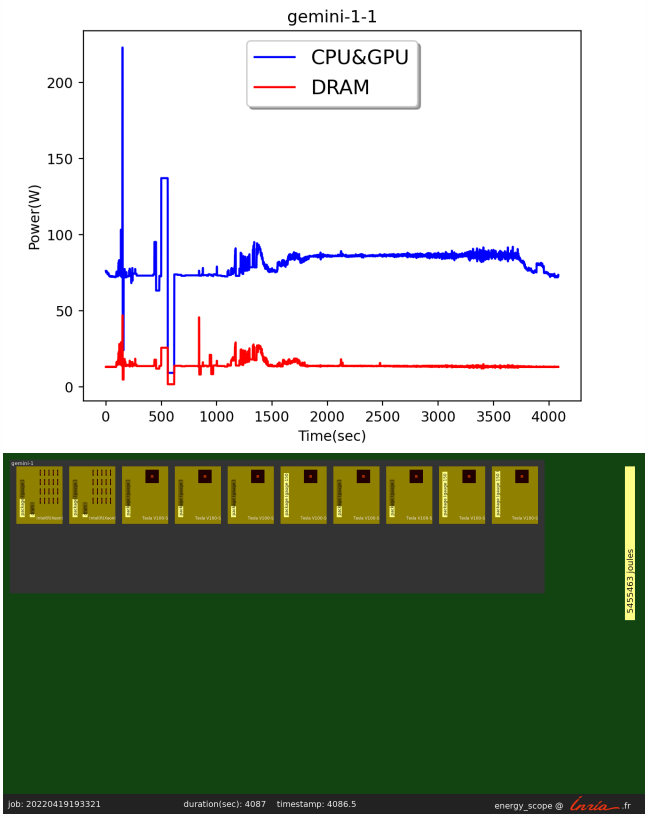

*node gemini-1/gpu-nvidia-0

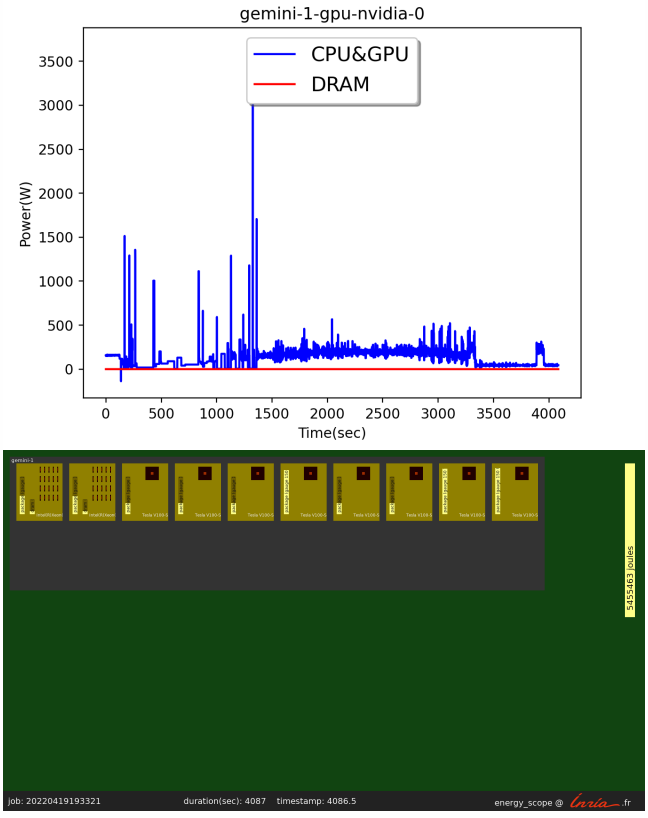

*node gemini-1/gpu-nvidia-1

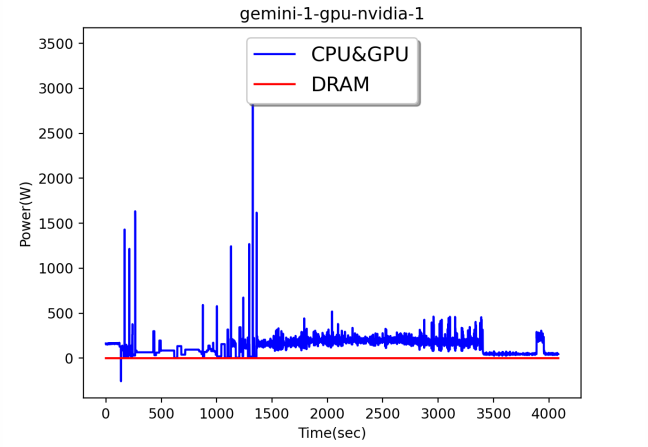

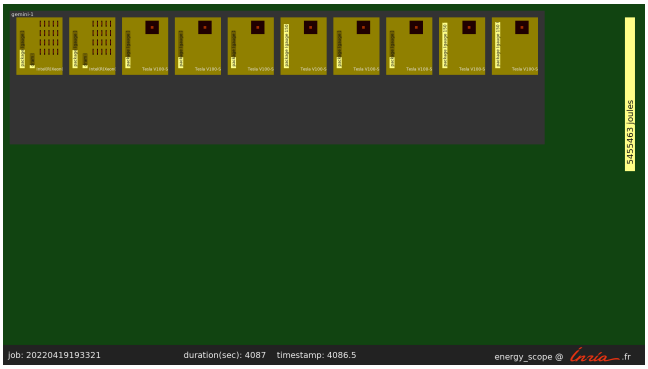

*node gemini-1/gpu-nvidia-2

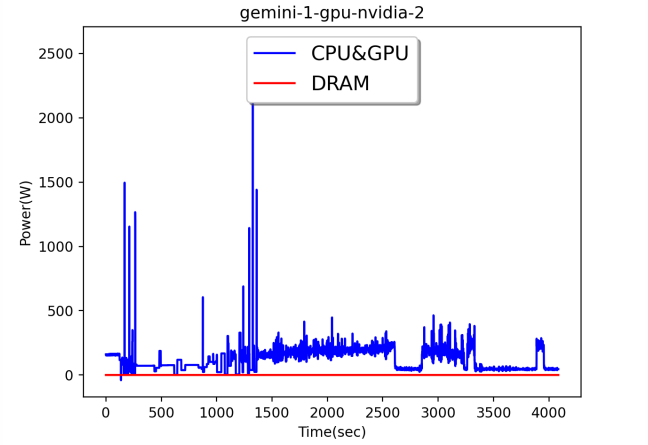

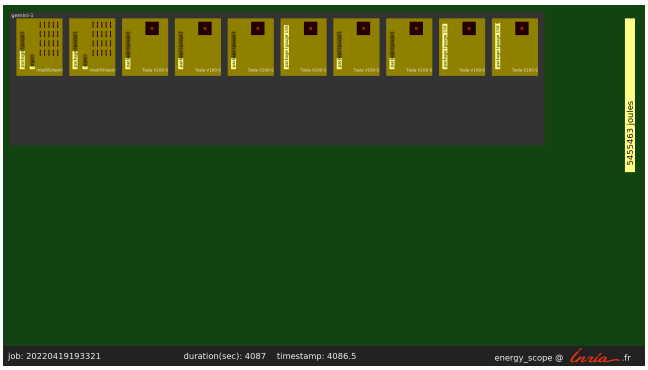

*node gemini-1/gpu-nvidia-3

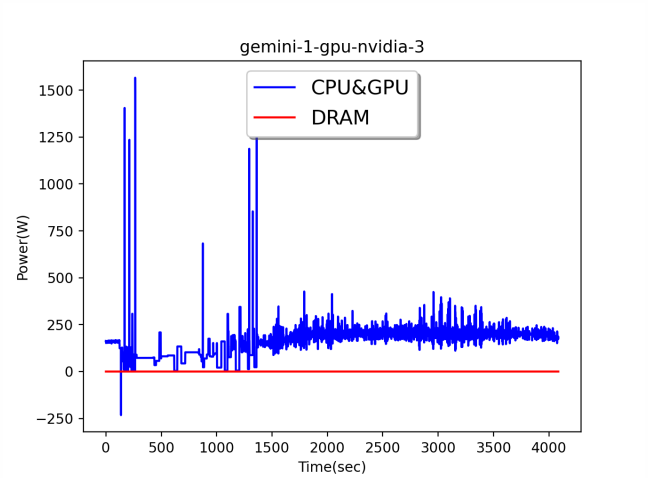

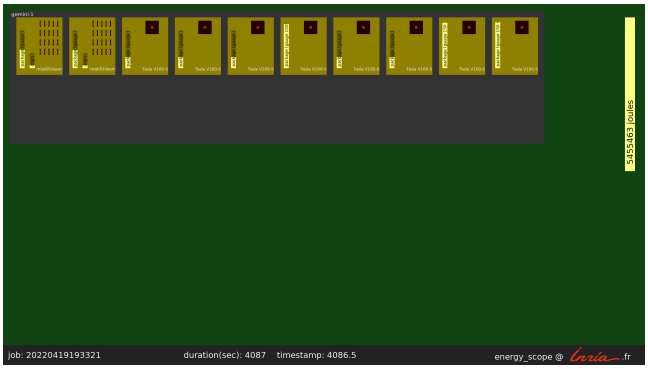

*node gemini-1/gpu-nvidia-4

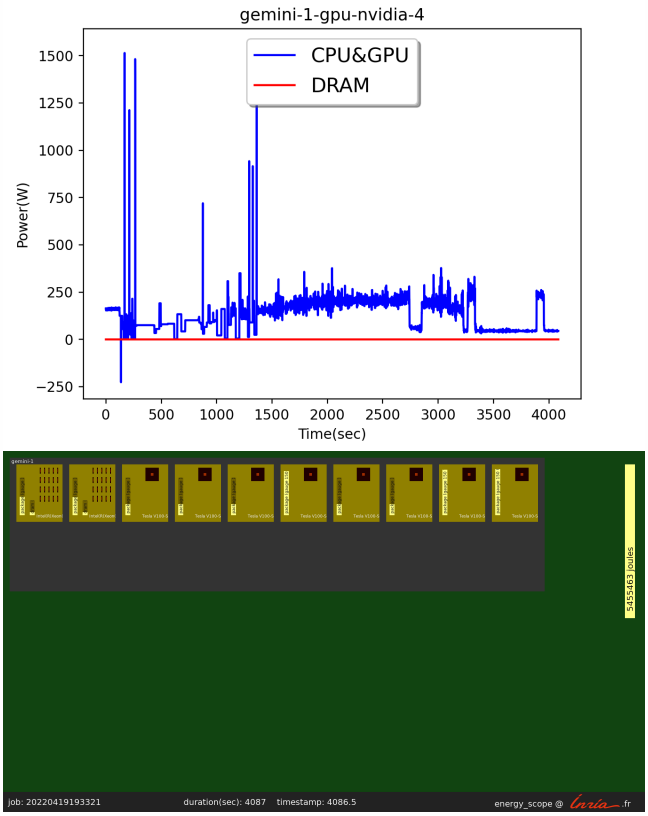

*node gemini-1/gpu-nvidia-5

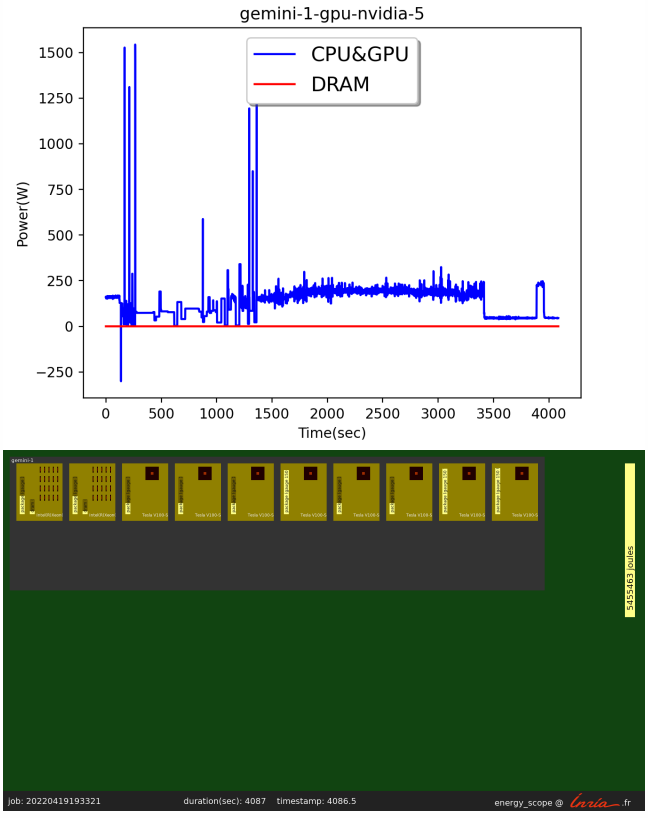

*node gemini-1/gpu-nvidia-6

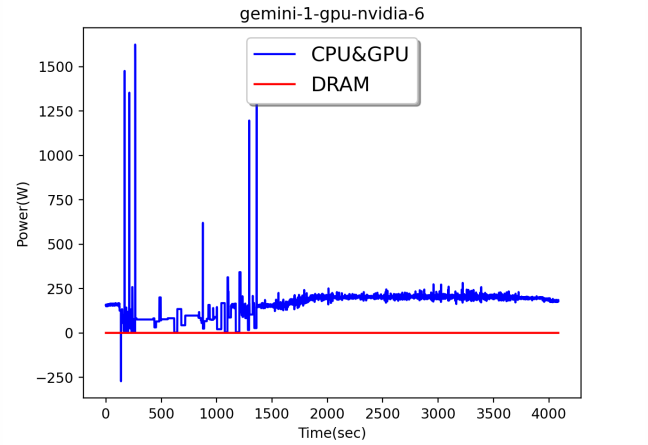

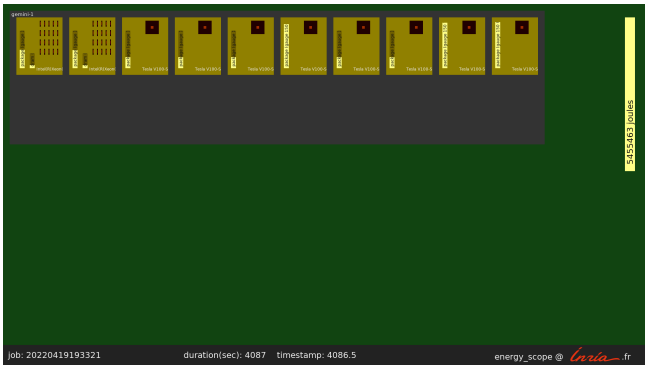

*node gemini-1/gpu-nvidia-7

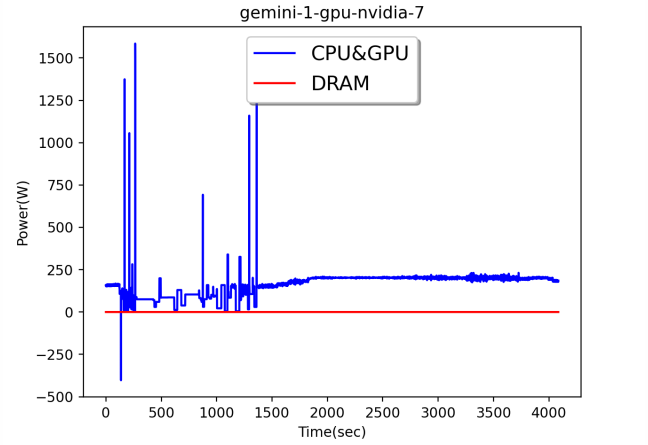

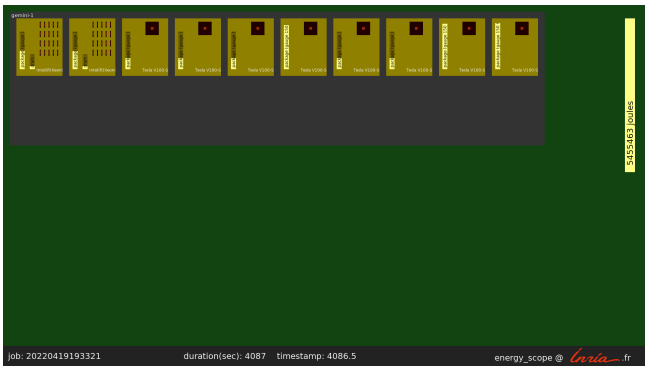
gemini-1
3455463 joules
job: 20220419193321
duration(sec): 4087
timestamp: 4086.5
energy_scope @ inria.fr