

Energy Scope report

Date of the report: 2022/04/19 18:37:20

GENERAL INFORMATION

- Jobid: 20220419193306
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_19_04/808_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/19 19:33:06.453516
- Duration (including ES prologue and epilogue): 3835 (sec)

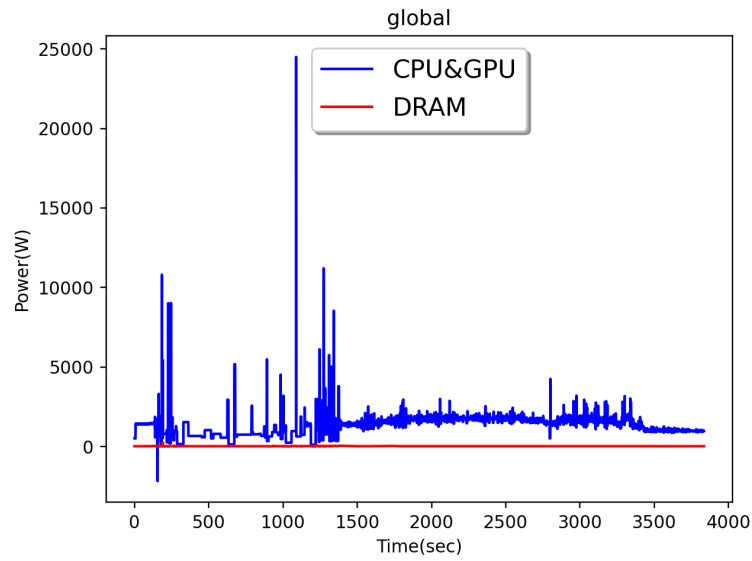
ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 1342.3 (J/sec)
- Application energy consumption measurement: 5147695 (J) 1.4299 (kWh)
- Global application energy consumption estimation: 6911795 (J) 1.9199 (kWh)
- Global application carbon production estimation (FR): 98.112 (gCO2)
- Energy efficiency (ref TDP): 59.13 (%)

Eprofile:



ENERGY ACQUISITION INFORMATION

- Period(ms): 1540.858
- Acquisition quality (low, medium, high): low
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	380130	63.2	59.1
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	367377	60.3	55.5
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	514461	53.7	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	532828	55.6	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	477262	49.8	20.0


node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	613721	64.0	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	512560	53.5	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	525850	54.8	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	610611	63.7	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	612895	63.9	20.0

PROFILES and CORE TEMPERATURE


Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.


*node gemini-1/0




*node gemini-1/1




*node gemini-1/gpu-nvidia-0

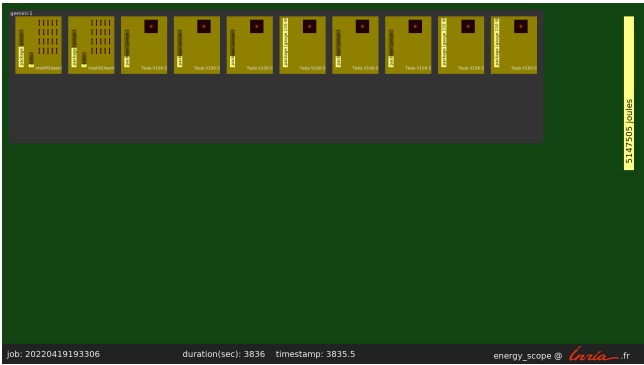


*node gemini-1/gpu-nvidia-1

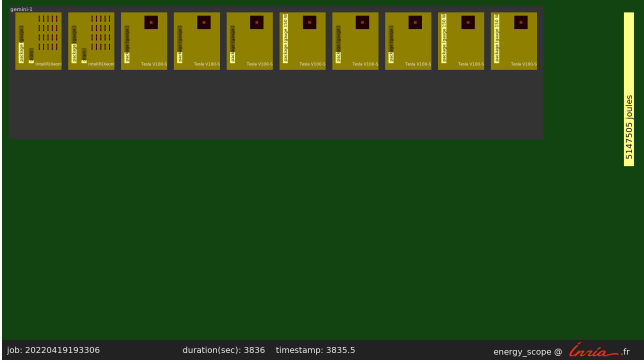



*node gemini-1/gpu-nvidia-2







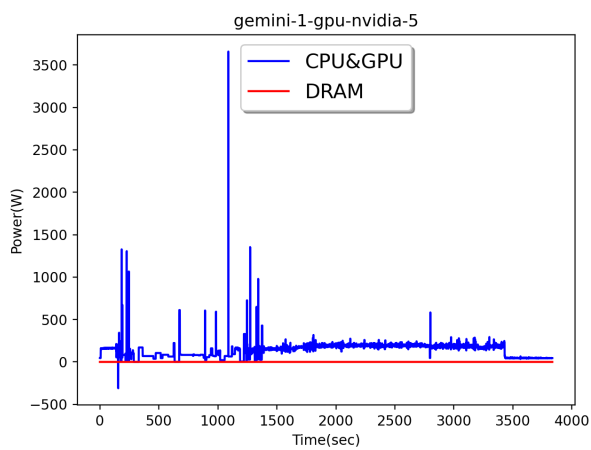
*node gemini-1/gpu-nvidia-3




*node gemini-1/gpu-nvidia-4



*node gemini-1/gpu-nvidia-5



*node gemini-1/gpu-nvidia-6



*node gemini-1/gpu-nvidia-7

