Energy Scope report

Date of the report: 2022/04/19 18:36:26

GENERAL INFORMATION

- Jobid: 20220419193317
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_19_04/936_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/19 19:33:17.866254
- Duration (including ES prologue and epilogue): 3770 (sec)

ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 1348.4 (J/sec)
- Application energy consumption measurement: 5083629 (J) 1.4121 (kWh)
- Global application energy consumption estimation: 6817829 (J) 1.8938 (kWh)
- Global application carbon production estimation (FR): 96.783 (gCO2)
- Energy efficiency (ref TDP): 59.4 (%)

Eprofile:

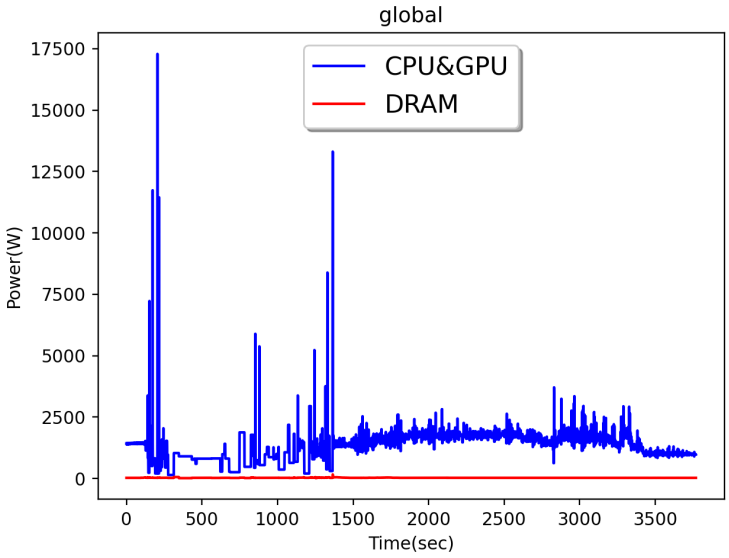

ENERGY ACQUISITION INFORMATION

- Period(ms): 1599.589
- Acquisition quality (low, medium, high): low
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	373904	63.3	59.1
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	361428	60.4	55.5
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	511071	54.2	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	529327	56.2	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	473774	50.3	20.0

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	602025	63.9	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	509118	54.0	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	522368	55.4	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	599147	63.6	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	601467	63.8	20.0

PROFILES and CORE TEMPERATURE

Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.

*node gemini-1/0

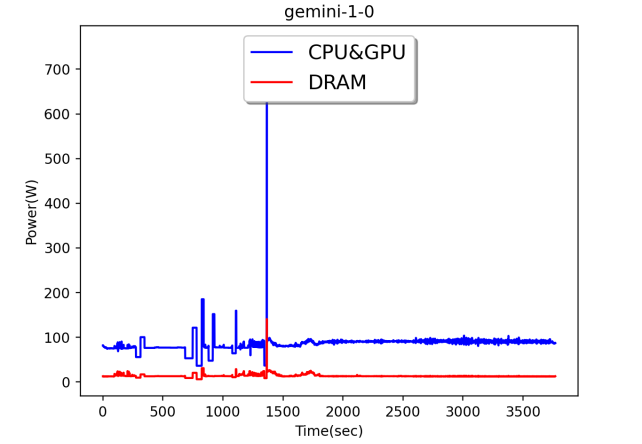

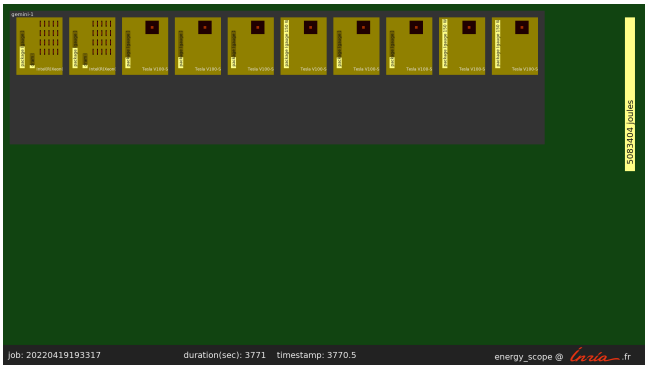

*node gemini-1/1

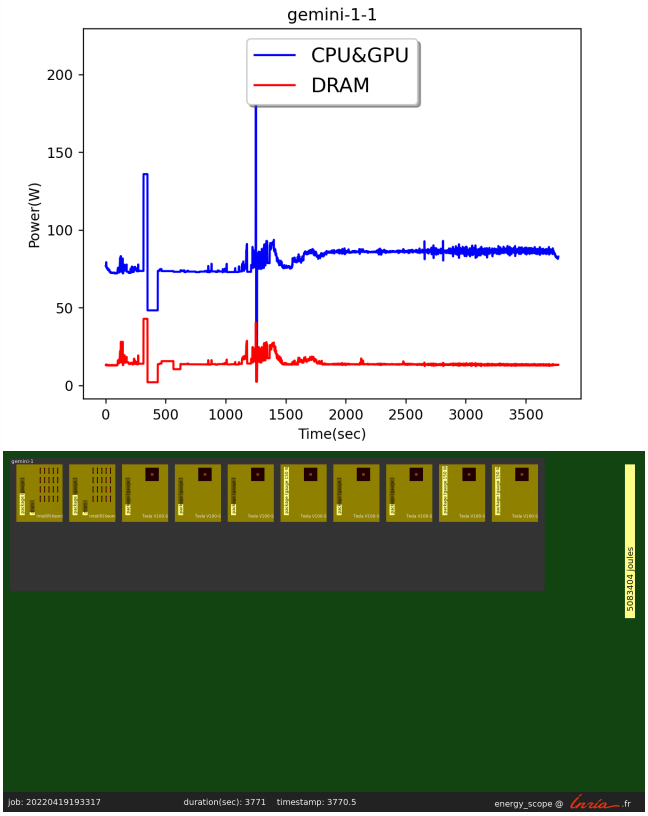

*node gemini-1/gpu-nvidia-0

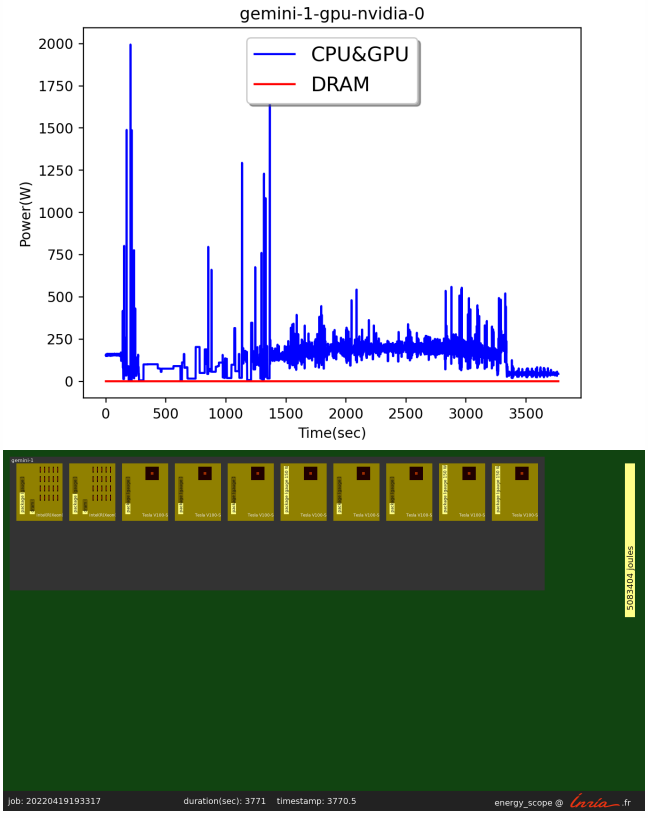

*node gemini-1/gpu-nvidia-1

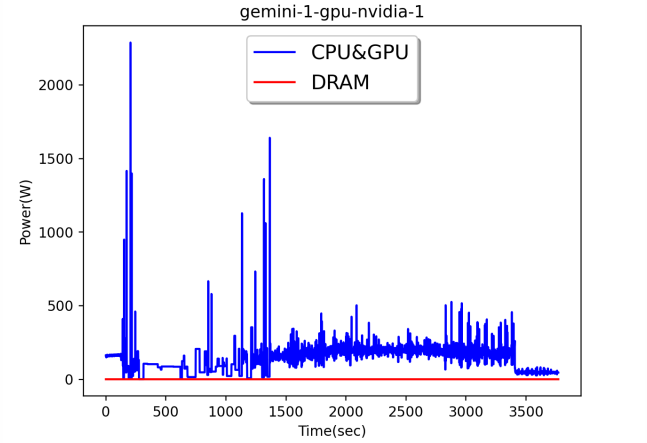

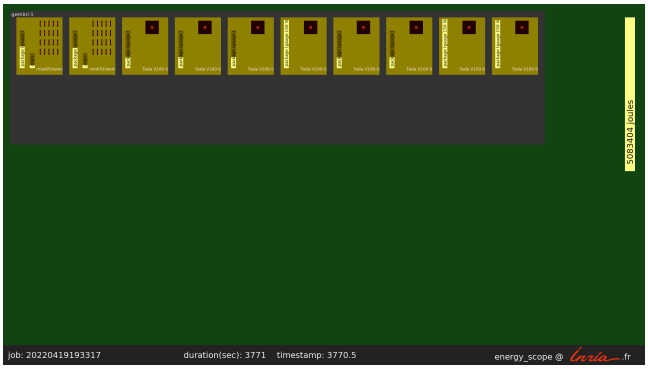

*node gemini-1/gpu-nvidia-2

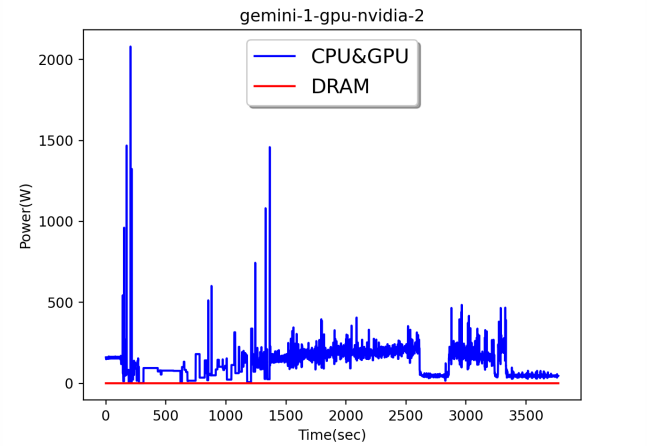

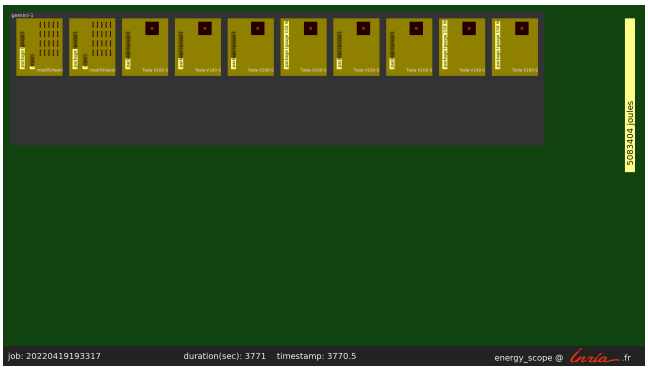

*node gemini-1/gpu-nvidia-3

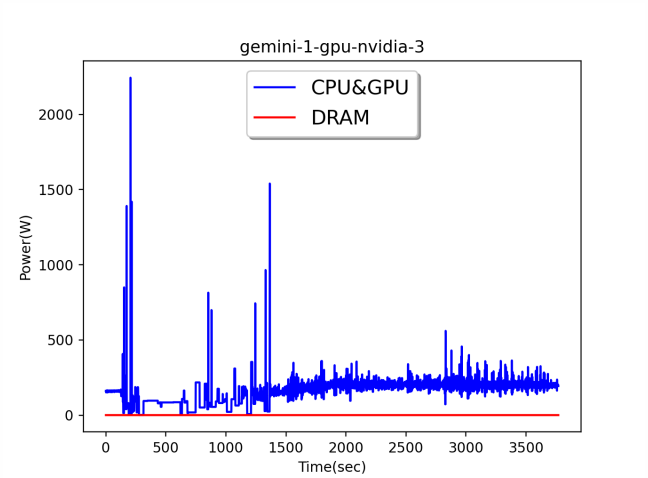

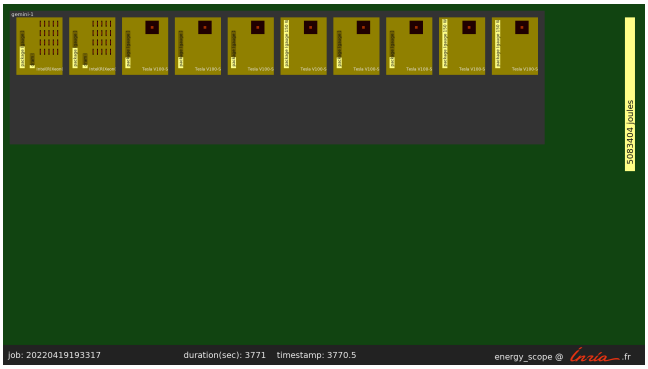

*node gemini-1/gpu-nvidia-4

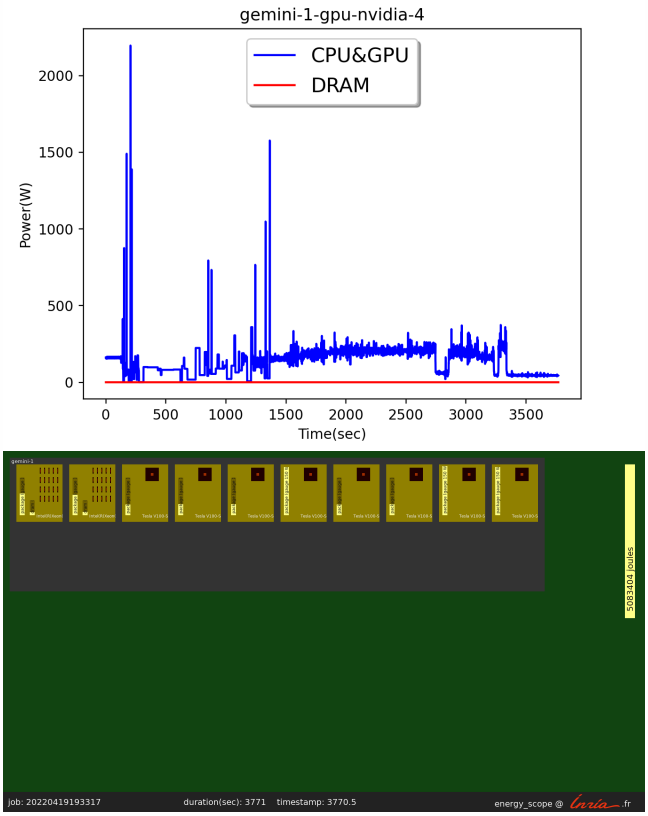

*node gemini-1/gpu-nvidia-5

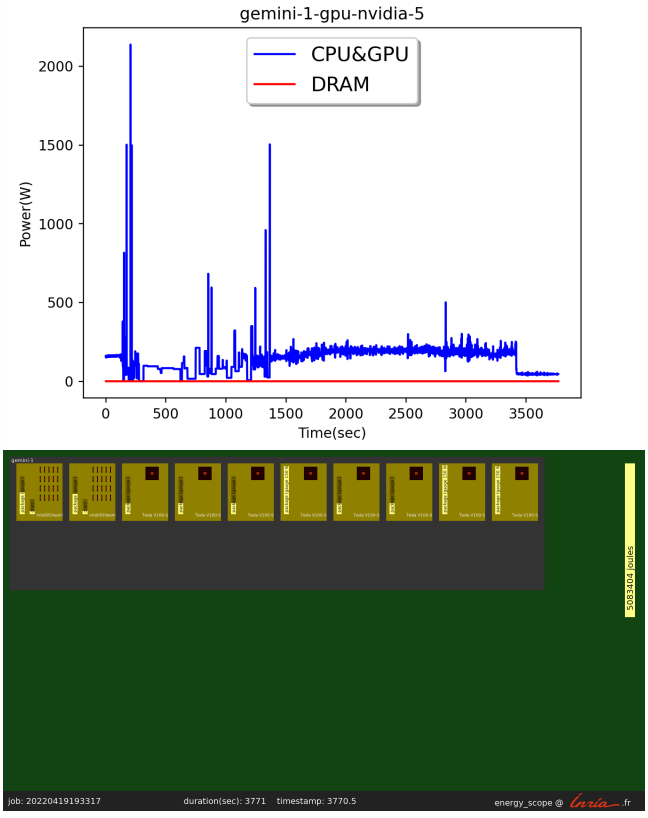

*node gemini-1/gpu-nvidia-6

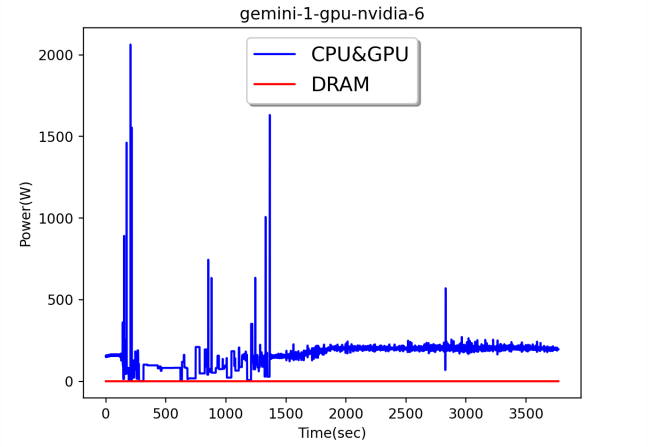

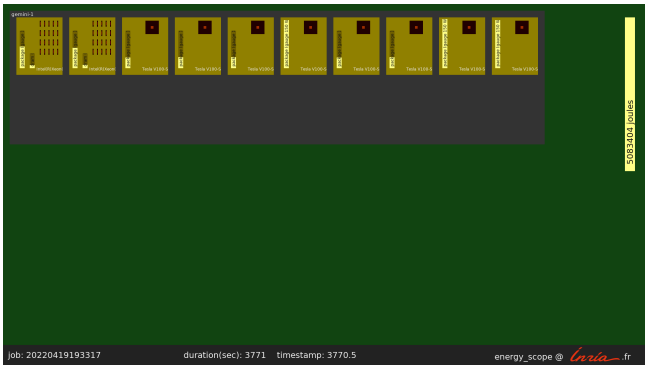

*node gemini-1/gpu-nvidia-7

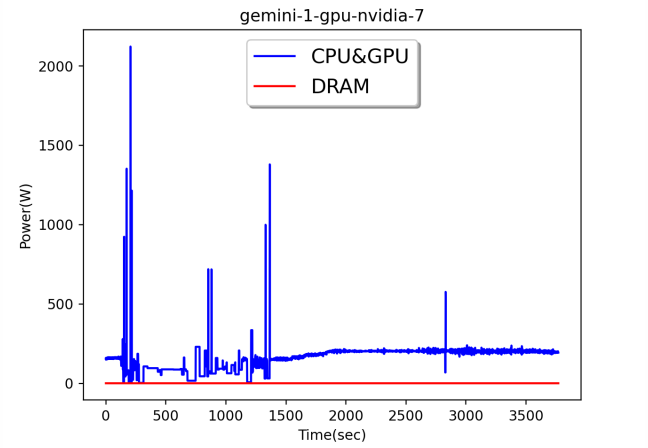

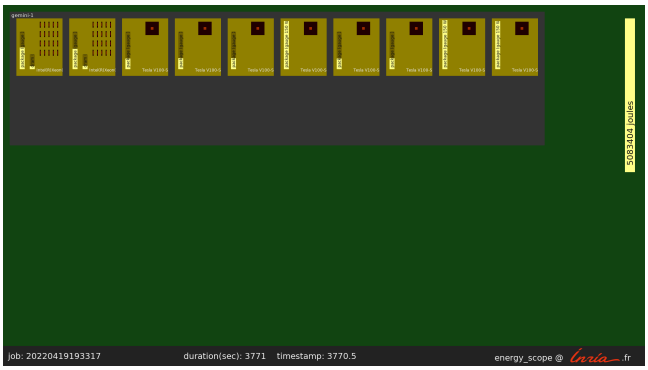
gemini-1
5083404 joules
job: 20220419193317
duration(sec): 3771
timestamp: 3770.5
energy_scope @ inria.fr