

Energy Scope report

Date of the report: 2022/04/20 19:04:56

GENERAL INFORMATION

- Jobid: 20220420205209
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_20_04/22_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/20 20:52:21.816254
- Duration (including ES prologue and epilogue): 738 (sec)


ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 965.6 (J/sec)
- Application energy consumption measurement: 712648 (J) 0.198 (kWh)
- Global application energy consumption estimation: 1052128 (J) 0.2923 (kWh)
- Global application carbon production estimation (FR): 14.921 (gCO2)
- Energy efficiency (ref TDP): 42.54 (%)

Eprofile:



ENERGY ACQUISITION INFORMATION

- Period(ms): 520.359
- Acquisition quality (low, medium, high): high
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	65540	55.8	58.3
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	63808	53.7	54.2
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	71776	38.9	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	74104	40.2	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	72146	39.1	20.0


node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	72745	39.4	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	73283	39.7	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	73277	39.7	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	73262	39.7	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	72707	39.4	20.0

PROFILES and CORE TEMPERATURE

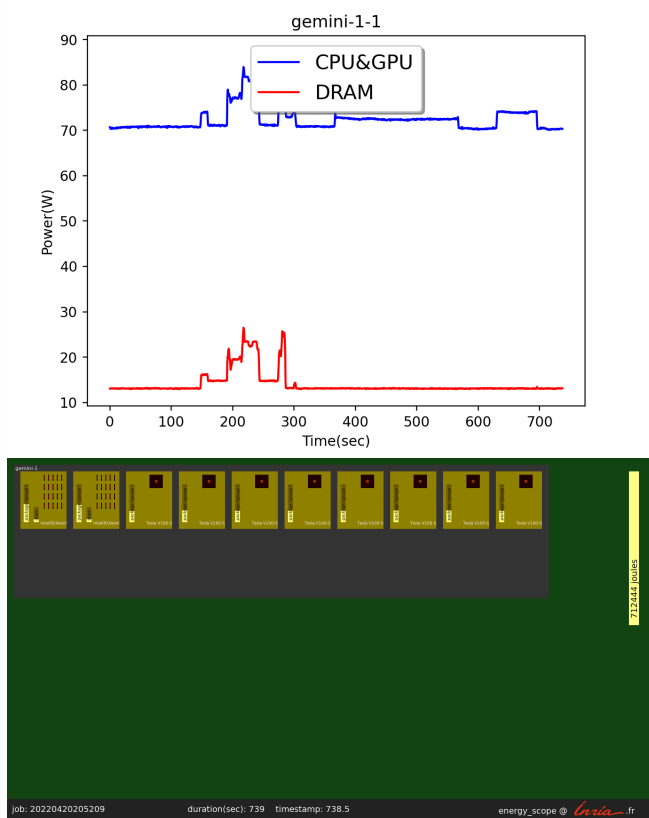
Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.


*node gemini-1/0




*node gemini-1/1



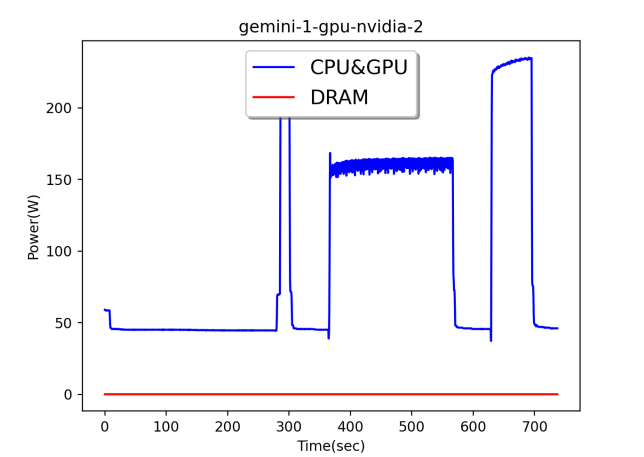
*node gemini-1/gpu-nvidia-0

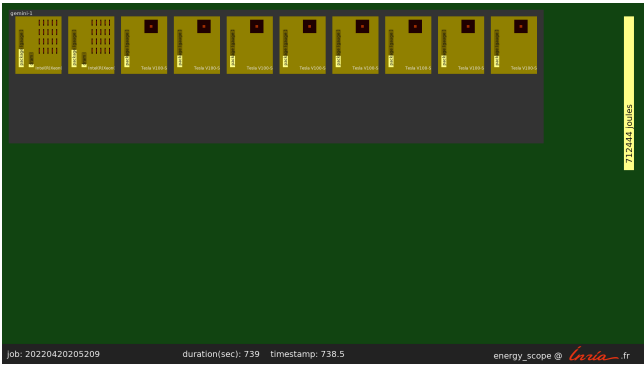


*node gemini-1/gpu-nvidia-1




*node gemini-1/gpu-nvidia-2







*node gemini-1/gpu-nvidia-3




*node gemini-1/gpu-nvidia-4



*node gemini-1/gpu-nvidia-5



*node gemini-1/gpu-nvidia-6



*node gemini-1/gpu-nvidia-7

