

Energy Scope report

Date of the report: 2022/04/20 18:25:02

GENERAL INFORMATION

- Jobid: 20220420201214
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_20_04/483_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/20 20:12:26.967643
- Duration (including ES prologue and epilogue): 738 (sec)


ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 969.8 (J/sec)
- Application energy consumption measurement: 715716 (J) 0.1988 (kWh)
- Global application energy consumption estimation: 1055196 (J) 0.2931 (kWh)
- Global application carbon production estimation (FR): 14.972 (gCO2)
- Energy efficiency (ref TDP): 42.72 (%)

Eprofile:



ENERGY ACQUISITION INFORMATION

- Period(ms): 520.181
- Acquisition quality (low, medium, high): high
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	65669	55.9	58.2
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	63849	53.7	54.1
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	71799	38.9	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	74083	40.2	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	72856	39.5	20.0


node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	73397	39.8	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	73984	40.1	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	73367	39.8	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	73944	40.1	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	72768	39.4	20.0

PROFILES and CORE TEMPERATURE


Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.


*node gemini-1/0



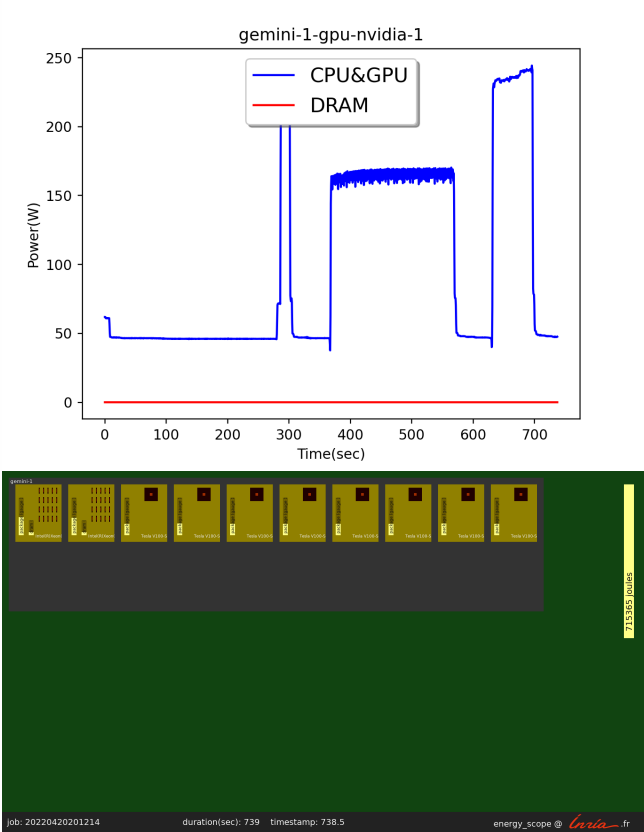
*node gemini-1/1



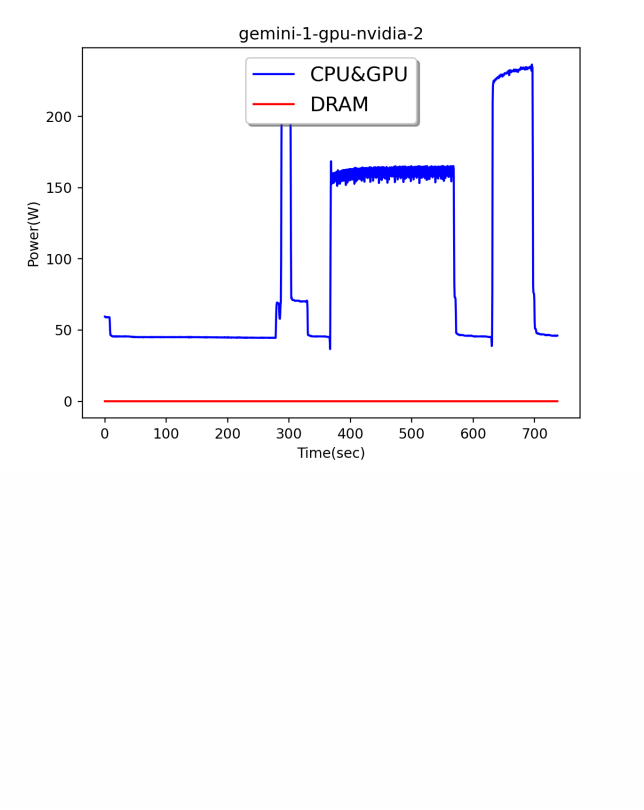
*node gemini-1/gpu-nvidia-0

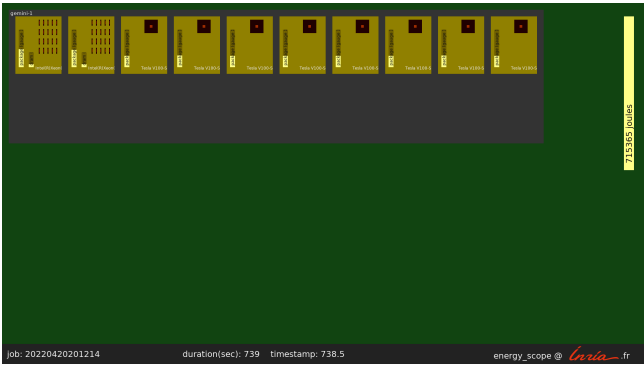


*node gemini-1/gpu-nvidia-1

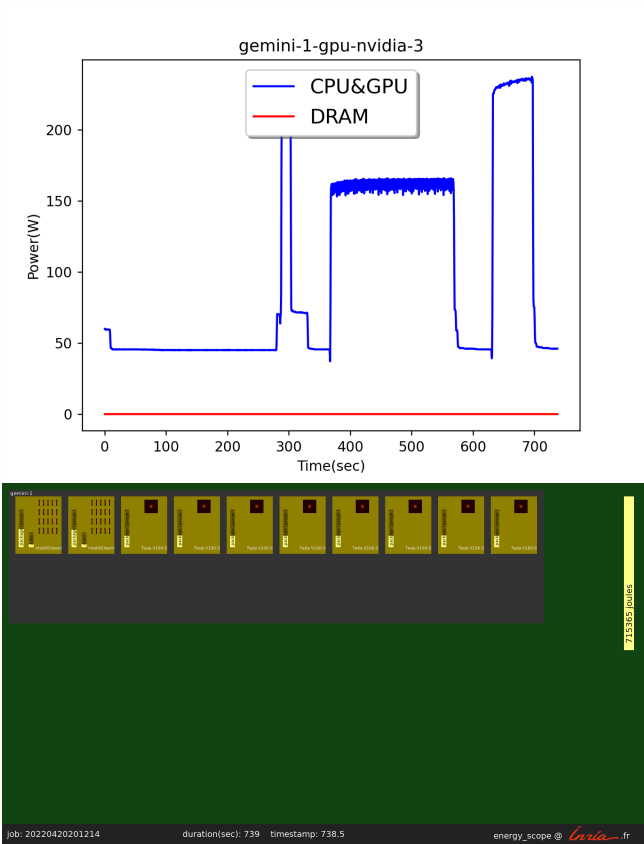


*node gemini-1/gpu-nvidia-2






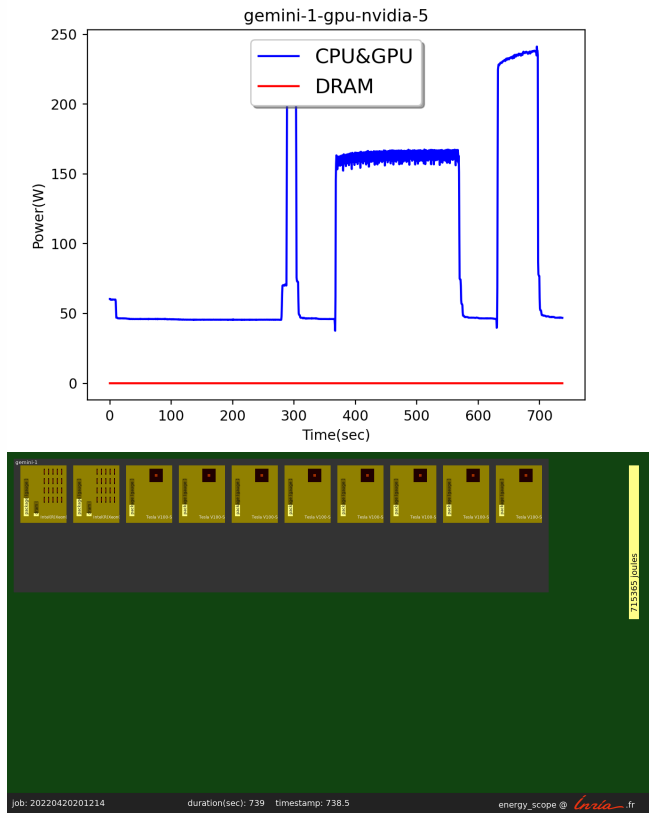
*node gemini-1/gpu-nvidia-3



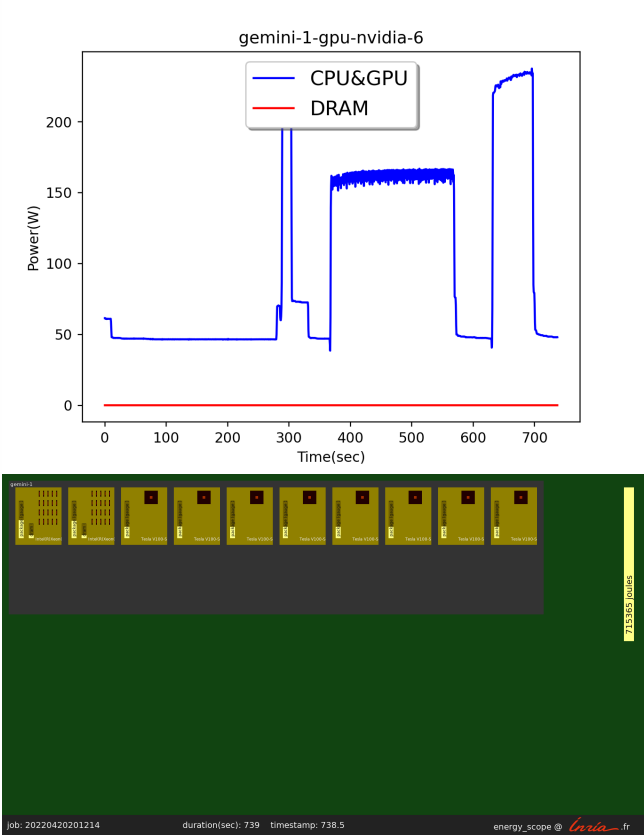
*node gemini-1/gpu-nvidia-4



*node gemini-1/gpu-nvidia-5



*node gemini-1/gpu-nvidia-6



*node gemini-1/gpu-nvidia-7

