

Energy Scope report

Date of the report: 2022/04/21 06:33:01

GENERAL INFORMATION

- Jobid: 20220421082012
- Command: /root/energy-consumption-of-gpu-benchmarks//results/night_exp_20_04/716_0//gpu0/scripts/script_final.sh
- Date of run: 2022/04/21 08:20:24.971917
- Duration (including ES prologue and epilogue): 738 (sec)


ARCHITECTURE INFORMATION

- nodelist: gemini-1
- processors type: Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz (TDP=135W)
- gpu type: Tesla V100-SXM2-32GB (TDP=250W)

ENERGY DATA

- Ratio Energy / Duration= 964.4 (J/sec)
- Application energy consumption measurement: 711698 (J) 0.1977 (kWh)
- Global application energy consumption estimation: 1051178 (J) 0.292 (kWh)
- Global application carbon production estimation (FR): 14.921 (gCO2)
- Energy efficiency (ref TDP): 42.48 (%)

Eprofile:



ENERGY ACQUISITION INFORMATION

- Period(ms): 520.12
- Acquisition quality (low, medium, high): high
- Information dumped: ecpu edram core_temperature

ENERGY BEHAVIOR

SUMMARY

node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
node gemini-1						
	cpu 0	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	66137	56.4	60.4
	cpu 1	Intel(R)Xeon(R)CPUE5-2698v4@2.20GHz	135	63917	53.8	54.4
	gpu gpu-nvidia-0	Tesla V100-SXM2-32GB	250	71528	38.8	20.0
	gpu gpu-nvidia-1	Tesla V100-SXM2-32GB	250	73717	40.0	20.0
	gpu gpu-nvidia-2	Tesla V100-SXM2-32GB	250	71853	38.9	20.0

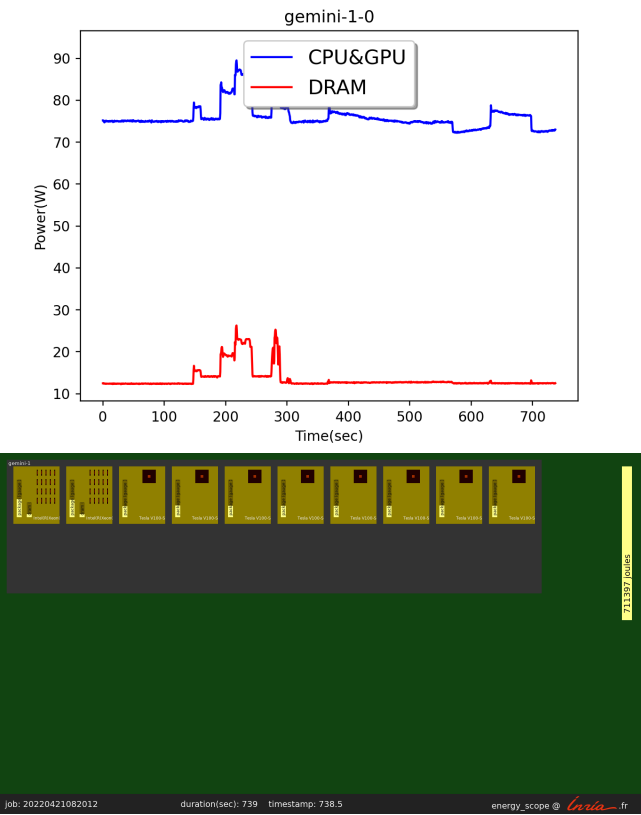
node	cpu/gpu	model	TDP (W)	Energy (J)	efficiency (%)	Cores Temp (C)
	gpu gpu-nvidia-3	Tesla V100-SXM2-32GB	250	72385	39.2	20.0
	gpu gpu-nvidia-4	Tesla V100-SXM2-32GB	250	73035	39.6	20.0
	gpu gpu-nvidia-5	Tesla V100-SXM2-32GB	250	73619	39.9	20.0
	gpu gpu-nvidia-6	Tesla V100-SXM2-32GB	250	73203	39.7	20.0
	gpu gpu-nvidia-7	Tesla V100-SXM2-32GB	250	72304	39.2	20.0

PROFILES and CORE TEMPERATURE

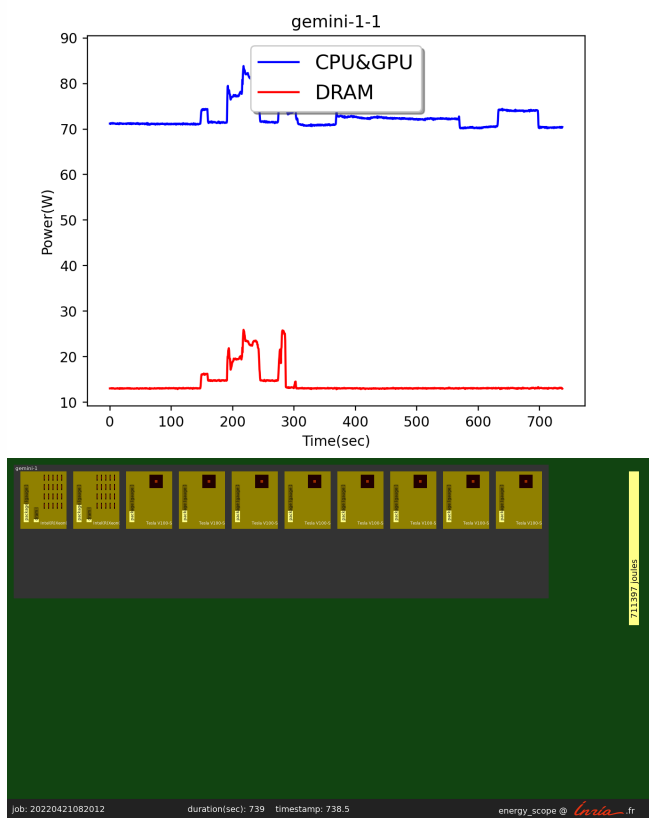
Images showing core temperature are generated when the average (of all the core) is maximum.

The full video showing the core temperature and the energy consumption over the time is available on demand.


*node gemini-1/0




*node gemini-1/1




*node gemini-1/gpu-nvidia-0

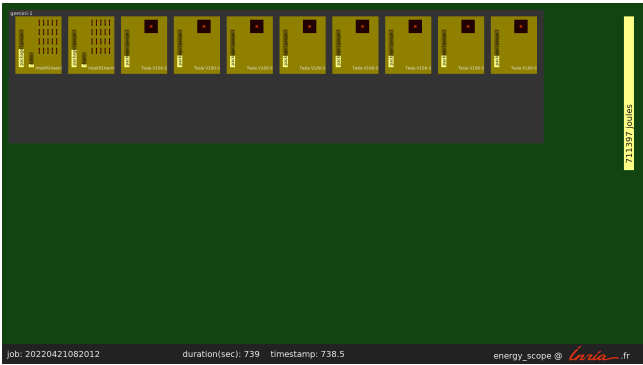


*node gemini-1/gpu-nvidia-1

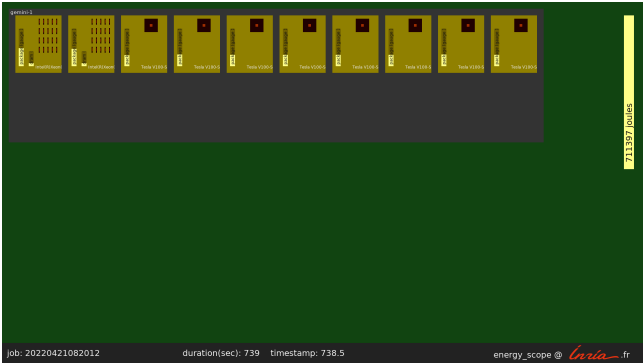
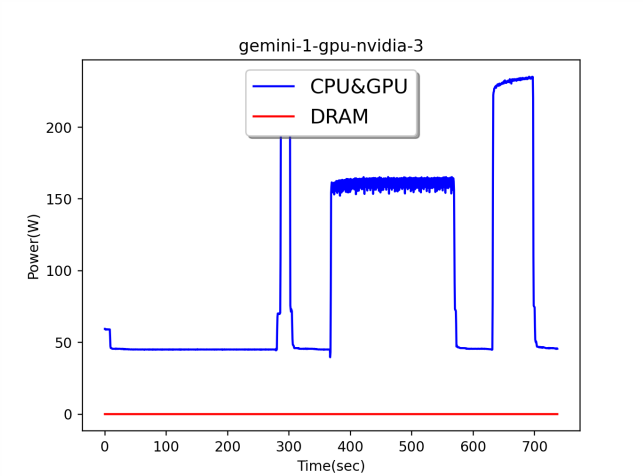


*node gemini-1/gpu-nvidia-2







*node gemini-1/gpu-nvidia-3




*node gemini-1/gpu-nvidia-4



*node gemini-1/gpu-nvidia-5



*node gemini-1/gpu-nvidia-6



*node gemini-1/gpu-nvidia-7

